

Module 3: Regression with multiple variables

Linear Regression with multiple variables

Multiple Features

Content

Multiple Linear Regression

- Multiple Features
- Vectorization
- Optional Lab: Python, NumPy and Vectorization
- Gradient Descent for Multiple Regression
- Optional Lab: Multiple Linear Regression
- Practical Quiz: Multiple Linear Regression

Gradient Descent in practice

- Feature Scaling
- Checking Gradient Descent for Convergence
- Choosing the Learning Rate
- Optional Lab: Feature Scaling and Learning rate
- Feature Engineering
- Polynomial Regression
- Optional Lab: Feature engineering and Polynomial Regression
- Optional Lab: Linear regression with Scikit-learn

What is this module is all about ...

- In this Module, you'll extend linear regression from single input to **multiple input features**.
- You'll also learn some methods for improving your model's training and performance, such as **vectorization**, **feature scaling**, **feature engineering** and **polynomial regression**.
- At the end of the module, you'll get to practice **implementing** linear regression using python.

Multiple Features (variables)

Multiple features (variables) is an extension of single feature regression

One
feature →

Size in feet ² (x)	Price (\$) in 1000's (y)	Output (Label)
2104	400	←
1416	232	
1534	315	
852	178	
...	...	

$$f_{w,b}(x) = wx + b$$

Uptill now we have discussed a regression model with only single input variabe, in this module we will introduce a regression model with multiple input variables

Multiple features (variables)

Size in feet ²	Number of bedrooms	Number of floors	Age of home in years	Price (\$) in \$1000's
x_1	x_2	x_3	x_4	
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

$i=2$

$j=1 \dots 4$

$n=4$

$x_j = j^{th}$ feature $j = 1, \dots, 4, n = 4$

n = number of features

$\vec{x}^{(i)}$ = features of i^{th} training example

$x^{(i)}$ = value of feature j in i^{th} training example

$$\vec{x}^{(2)} = [1416 \ 3 \ 2 \ 40]$$

$x_3^{(2)} = 2$ this is 3rd value of the 2nd training example, of the 3rd feature X3

Model with Multiple Features

Previously: $f_{w,b}(x) = wx + b$ single feature x

General Model

$$f_{w,b}(x) = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$$

Example

$$f_{w,b}(x) = 0.1 \underset{\text{size}}{\uparrow} x_1 + 4 \underset{\text{\#bedrooms}}{\uparrow} x_2 + 10 \underset{\text{\#floors}}{\uparrow} x_3 + -2 \underset{\text{years}}{\uparrow} x_4 + 80 \underset{\text{base price}}{\uparrow}$$

$$f_{w,b}(x) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

How to interpret coefficients

$$Y = 80 + w_1 \cdot X_1 + w_2 \cdot X_2 + w_3 \cdot X_3 + w_4 \cdot X_4$$

Estimated Price = $80 + w_1 \cdot \text{Size} + w_2 \cdot \text{Bedrooms} + w_3 \cdot \text{Floors} + w_4 \cdot \text{Age}$

Estimated Price = $80 + 0.1 \cdot \text{Size} + 4 \cdot \text{Bedrooms} + 10 \cdot \text{Floors} + -2 \cdot \text{Age}$

- Size: The unit for the coefficient 0.1 would be \$100 (dollars) per square foot.
- Number of Bedrooms: The unit for the coefficient 4 would be 4×1000 (dollars) per additional bedroom.
- Number of Floors: The unit for the coefficient 10 would be $10 \times \$1000$ (dollars) per additional floor.
- House Age : The unit for the coefficient -2 would be $-2 \times \$1000$ (dollars) per additional year of age.

Question:

In the training set below, what is $x_1^{(4)}$, $x_3^{(3)}$, $x_5^{(1)}$, $x_2^{(2)}$

Size in feet ²	Number of bedrooms	Number of floors	Age of home in years	Price (\$) in \$1000's
x_1	x_2	x_3	x_4	
2104	5	1	45	460
1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

Multiple linear regression model expression (Vector Notation)

$$f_{w,b}(x) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

$$\vec{w} = [w_1 \ w_2 \ w_3 \ \dots \ w_n] \text{ (Vector)}$$

b is a number

parameters
of the model

vector $\vec{x} = [x_1 \ x_2 \ x_3 \ \dots \ x_n]$

So the model can be written as

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b = w_1x_1 + w_2x_2 + w_3x_3 + \dots + w_nx_n + b$$

dot product

Multiple Linear Regression

Vectorization Part 1

What is Vectorization

- Vectorization in multivariate regression involves using mathematical operations on entire sets of data at once, taking advantage of linear algebra principles. Instead of individually processing each data point and feature in loops.
- Vectorization allows for efficient computations by operating on matrices as a whole.
- This approach improves computational speed and is particularly useful for handling large datasets in machine learning, making the code more concise and leveraging optimized, hardware-accelerated operations on matrices.

Benefits of vectorization

Benefit	Description
Computational Efficiency	Vectorization significantly reduces computation time by processing multiple data points simultaneously.
Scalability	Vectorized algorithms can handle large datasets effectively, making them scalable to real-world applications.
Simplicity	Vectorized code is often more concise and easier to understand compared to iterative code.

Parameters and features with vectorization

for n = 3 $\vec{w} = [w_1 \quad w_2 \quad w_3]$ $\vec{x} = [x_1 \quad x_2 \quad x_3]$ b is a number

linear algebra: count from 1

```
w[0]  w[1]  w[2]
w = np.array([1.0, 2.5, -3.3])
b = 4
x[0]  x[1]  x[2]
x = np.array([10, 20, 30])
```

code: count from 0

NumPy  Python Linear algebra library

Without vectorization vs. with Vectorization

no vectorization Direct multiplication

$$f_{\vec{w}, b}(\vec{x}) = w_1 x_1 + w_2 x_2 + w_3 x_3 + b$$

```
f = w[0] * x[0] +
    w[1] * x[1] +
    w[2] * x[2] + b
```

np library uses parallel processing

No vectorization For loop

$$f_{\vec{w}, b}(\vec{x}) = \left(\sum_{j=1}^n w_j x_j \right) + b \quad \sum_{j=1}^n \rightarrow j=1 \dots n$$

range (0,n) = 0, 1, 2, 3, ... n-1

```
f = 0
for j in range(0, n):
    f = f + w[j] *
        x[j]
f = f + b
```

with Vectorization

$$f_{\vec{w}, b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$$

```
f = np.dot(w, x) + b
```

Question

Which of the following is a vectorized implementation for computing a linear regression model's prediction?

- `f = np.dot(w, x) + b`
- `f=0`
 For `j` in `range(n)` :
 `f = f + w[j] * x[j]`
`f=f + b`

To conclude ...

Vectorization is a critical technique that makes your code simpler, easier to write and to read and it runs much faster

But what is the magic behind vectorization that makes it run so much faster ???

lets see in the next section what computers do behind the scenes to make computers run much faster



Implementing Vectorization

The secret that makes vectorization make computers run faster

Without vectorization for loop

```
for j in range(0,16):  
    f = f + w[j] * x[j]
```

t_0 (operates on the 1st term

$f + w[0] * x[0]$

t_1

$f + w[1] * x[1]$

...


t_{15}

$f + w[15] * x[15]$

Vectorization

```
np.dot(w,x)
```

t_0 (process done at one time



Vectorization: efficient calculation and scale to large data set

Gradient descent with vectorization

$$\vec{w} = (w_1 \quad w_2 \quad \cdots \quad w_{16}) \quad \cancel{b} \text{ parameters}$$

$$\text{Derivatives } \vec{d} = (d_1 \quad d_2 \quad \cdots \quad d_{16})$$

Store values
in arrays

```
w = np.array([0.5, 1.3, ... 3.4])  
d = np.array([0.3, 0.2, ... 0.4])
```

coefficients
derivatives

compute (update)

$$w_j = w_j - 0.1 d_j \quad \text{for } j = 1 \dots 16$$

↑
learning rate

update

Without vectorization

$$w_1 = w_1 - 0.1 d_1$$

$$w_2 = w_2 - 0.1 d_2$$

⋮

$$w_{16} = w_{16} - 0.1 d_{16}$$


for j in range(0,16):

$$w[j] = w[j] - 0.1 * d[j]$$

With vectorization

$$\vec{w} = \vec{w} - 0.1 \vec{d}$$

do the calculation
at the same time



$$w = w - 0.1 * d$$

Question

Which of the following is a vectorized implementation for computing a linear regression model's prediction?


```
f = np.dot(w,x) + b
```

```
f = w[0] * x[0] + w[1] * x[1] + w[2] * x[2] + b
```

```
f = 0  
for j in range(n):  
    f = f + w[j] * x[j]  
f = f + b
```

To conclude ...

- You've learned one of the most important and useful techniques in implementing machine learning algorithms.
- In the next section, we'll put the math of multiple linear regression together with vectorization, so that you will implement gradient descent for multiple linear regression with vectorization.
- Let's go on to the next section.



Gradient Descent for Multiple Regression With Vectorization

You've learned about gradient descents about multiple linear regression and also vectorization.

Let's put it all together to implement gradient descent for multiple linear regression with vectorization.

Previous notation

Parameters

$$\begin{array}{c} w_1, \dots, w_n \\ b \end{array} \quad \longrightarrow \quad \longrightarrow$$

Model $f_{\vec{w}, b}(\vec{x}) = w_1 x_1 + \dots + w_n x_n + b$

Cost function $J(w_1, \dots, w_n, b)$

Gradient descent

repeat {

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(w_1, \dots, w_n, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w_1, \dots, w_n, b)$$

}

Vector notation

\vec{w} *vector of length n*
 $w_1 \dots w_n$
 b still a number
 $f_{\vec{w}, b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$

$$J(\vec{w}, b)$$

repeat {

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(\vec{w}, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(\vec{w}, b)$$

}

Gradient descent

One feature

repeat {

$$w = w - \alpha \left(\frac{1}{m} \sum_{i=1}^m f(w, b)(x^{(i)}) - y^{(i)} \right) x^{(i)}$$

 $\frac{\partial}{\partial w} J(w, b)$

$$b = b - \alpha \left(\frac{1}{m} \sum_{i=1}^m f(w, b)(x^{(i)}) - y^{(i)} \right)$$

simultaneously update w, b

}

n features ($n \geq 2$)

repeat {

j = 1

$$w_1 = w_1 - \alpha \left(\frac{1}{m} \sum_{i=1}^m f(w, b)(x^{(i)}) - y^{(i)} \right) x^{(i)}$$

.

 $\frac{\partial}{\partial w_1} J(w, b)$

$$w_n = w_n - \alpha \left(\frac{1}{m} \sum_{i=1}^m f(w, b)(x^{(i)}) - y^{(i)} \right) x^{(i)}$$

$$b = b - \alpha \left(\frac{1}{m} \sum_{i=1}^m f(w, b)(x^{(i)}) - y^{(i)} \right)$$

}

simultaneous update
 w_j for $j = 1 \dots n$ and b

An alternative to gradient descent

Normal equation

- Only for linear regression
- Solve for w , b without iterations

Disadvantages

- Doesn't generalize to other learning algorithms.
- Slow when number of features is large ($> 10,000$)

What you need to know

- Normal equation method may be used in machine learning libraries that implement linear regression.
- Gradient descent is the recommended method for finding parameters w, b

To conclude ...

- You now know multiple linear regression.
- This is probably the single most widely used learning algorithm in the world today.
- But there's more. With just a few tricks such as picking and scaling features appropriately and also choosing the learning rate alpha appropriately, you'd really make this work much better.
- Let's go on to the next section to see those little tricks that will help you make multiple linear regression work much better.

Practical Tips for Linear Regression

Feature Scaling

Simple definiton of feature scaling

- Imagine you have a dataset with two features: age and income. The age feature might range from 10 to 100, while the income feature might range from \$10,000 to \$1,000,000. When you train a machine learning model on this data, the model will give more weight to the income feature because it has a larger scale. This can make it difficult for the model to learn the relationship between the age feature and the target variable.
- Feature scaling is a technique that can be used to normalize the features in a dataset so that they all have a similar scale. This can help to improve the performance of machine learning models.

Feature and parameter values

$$price = w_1 x_1 + w_2 x_2 + b$$



x_1 : size (feet²)
range: 300 – 2,000

x_2 : # bedrooms
range: 0 – 5
small

House: $x_1 = 2000$, $x_2 = 5$, $price = \$500k$ one training example

size of the parameters w_1, w_2 ?

$$w_1 = 50, \quad w_2 = 0.1, \quad b = 50$$

large small

$$w_1 = 0.1, \quad w_2 = 50, \quad b = 50$$

small large

$$price = 50 * 2000 + 0.1 * 5 + 50$$

100,000 k 0.5 k 50 k

$$price = 0.1 * 2000k + 50 * 5 + 50$$

200 k 250 k 50 k

$$price = \$100,050.5k$$

$$price = \$500k$$

More Reasonable


What do we conclude

- So hopefully you might notice that when a possible range of values of a feature is large, like the size and square feet which goes all the way up to 2000. It's more likely that a good model will learn to choose a relatively small parameter value, like 0.1.
- Likewise, when the possible values of the feature are small, like the number of bedrooms, then a reasonable value for its parameters will be relatively large like 50.



So how is that relate to gradient Descent ...?

Feature size and parameter size

	size of feature x_j	size of parameter w_j
size in feet ²	↔	↔
#bedrooms	↔	↔



Feature size and gradient descent




The rescaled values of the parameters are taking comparable ranges to each others

Practical Tips for Linear Regression

How to implement Feature scaling

Feature scaling



$$300 \leq x_1 \leq 2000$$

$$x_{1,scaled} = \frac{x_1}{\max}$$



$$0.15 \leq x_{1,scaled} \leq 1$$

$$0 \leq x_2 \leq 5$$

$$x_{2,scaled} = \frac{x_2}{\max}$$

$$0 \leq x_{2,scaled} \leq 1$$

Mean normalization



$$300 \leq x_1 \leq 2000$$

$$x_1 = \frac{x_1 - \mu_1}{2000 - 300}$$

max-min

$$-0.18 \leq x_1 \leq 0.82$$



$$0 \leq x_2 \leq 5$$

$$x_2 = \frac{x_2 - \mu_2}{5 - 0}$$

max-min

$$-0.46 \leq x_2 \leq 0.54$$

Z-score normalization



$$300 \leq x_1 \leq 2000$$

$$x_1 = \frac{x_1 - \mu_1}{\sigma_1}$$

$$-0.67 \leq x_1 \leq 3.1$$

$$0 \leq x_2 \leq 5$$

$$x_2 = \frac{x_2 - \mu_2}{\sigma_2}$$

$$-1.6 \leq x_2 \leq 1.9$$

Feature scaling

aim for about $-1 \leq x_j \leq 1$ for each feature x_j

$$\begin{array}{c} -3 \leq x_j \leq 3 \\ -0.3 \leq x_j \leq 0.3 \end{array} \quad \left. \begin{array}{l} \\ \end{array} \right\} \text{acceptable ranges}$$

$$0 \leq x_1 \leq 3$$

Okay, no rescaling

$$-2 \leq x_2 \leq 0.5$$

Okay, no rescaling

$$-100 \leq x_3 \leq 100$$

too large \rightarrow rescale

$$-0.001 \leq x_4 \leq 0.001$$

too small \rightarrow rescale


$$98.6 \leq x_5 \leq 105$$

too large \rightarrow rescale

To conclude ...

- There's almost never any harm to carrying out feature re-scaling. When in doubt, I encourage you to just carry it out.
- That's it for feature scaling. With this little technique, you'll often be able to get gradient descent to run much faster. That's features scaling.
- With or without feature scaling, when you run gradient descent, how can you know, how can you check if gradient descent is really working?
- If it is finding you the global minimum or something close to it.

In the next section, let's take a look at how to recognize if gradient descent is converging, and then in the video after that, this will lead to discussion of how to choose a good learning rate for gradient descent.



Checking Gradient Descent for Convergence

Gradient descent

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(w_1, \dots, w_n, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w_1, \dots, w_n, b)$$

these are the gradient descent algorithm equations, and one of the main choices of the parameters is the choice of alpha (a) the learning rate

So the question is how do we choose alpha

- 1.The number of iterations in gradient descent significantly impacts the cost function.
- 2.As the number of iterations increases, the cost function generally decreases.
- 3.The optimal number of iterations depends on various factors and may not be linear.

Make sure gradient descent is working correctly


objective: $\min_{\vec{w}, b} J(\vec{w}, b)$

$J(\vec{w}, b)$ should decrease after every iteration


plot $J(w, b)$ as a function of number of iterations.

→ the cost function should decrease constantly

if the cost function increase after some iterations, then either alpha is chosen poorly, or there is a bug in the code



Make sure gradient descent is working correctly





- Number of iterations varies for different applications, it may take 30, some other applications might take 1000, or 10,000 iterations.
- it might be difficult to tell in advance how many iterations gradient descent might need to converge

to conclude ...

- You've now seen what the learning curve should look like when gradient descent is running well.
- Let's take these insights and in the next part, take a look at how to choose an appropriate learning rate.

Choosing the Learning Rate

Identify problem with gradient descent





$w_1 = w_1 + \alpha d_1$ 
use a minus sign
 $w_1 = w_1 - \alpha d_1$ 



or

learning rate is too large

Identify problem with gradient descent



With a small enough α ,
 $J(\vec{w}, b)$ should decrease
on every iteration





If α is too small,
gradient descent takes
a lot more iterations to
converge

Since alpha is big you may skip the minimum,
and that will make the cost function increases

Values of α to try:

... 0.001 0.003 0.01 0.03 0.1 0.3 1 ...
3X $\approx 3X$ 3X $\approx 3X$ 3X $\approx 3X$



Feature Engineering

Feature engineering is a crucial step in machine learning that involves transforming raw data into features suitable for machine learning algorithms.

Key objectives of feature engineering:

- Select relevant features: Identify and select the most relevant features from the dataset to eliminate noise and focus on important factors.
- Create new features: Derive new features from existing data to capture underlying relationships and patterns that may not be directly present in the raw data.
- Transform existing features: Transform existing features using techniques like scaling, normalization, and binning to improve model convergence and stabilize the training process.

Benefits of feature engineering:

- Improved model performance: Enhances the accuracy and generalizability of machine learning models.
- Faster model training: Reduces the complexity of data and improves model convergence.
- Better understanding of data: Reveals underlying relationships and patterns in the data, leading to more meaningful insights.

Feature engineering

$$f_{\vec{w}, b}(\vec{x}) = w_1 x_1 + w_2 x_2 + b$$


frontage depth

$$\text{area} = \text{frontage} \times \text{depth}$$

$$x_3 = x_1 x_2$$

new feature

$$f_{\vec{w}, b}(\vec{x}) = w_1 x_1 + w_2 x_2 + w_3 x_3 + b$$



Feature engineering: Using intuition to design new features, by transforming or combining original features.


So by designing a new feature you might get a better model

Polynomial Regression

Polynomial regression is a type of regression analysis that models the relationship between a response variable and one or more independent variables using a polynomial equation.

In other words, it assumes that the relationship between the variables can be expressed as a sum of terms, where each term is the product of a constant and a variable raised to an integer power.

Polynomial regression



$$f_{\vec{w},b}(x) = w_1x + w_2x^2 + w_3x^3 + b$$

$$\begin{array}{c} \uparrow \\ \text{size} \\ \underline{1-10^3} \end{array} \quad \begin{array}{c} \uparrow \\ \text{size}^2 \\ \underline{1-10^6} \end{array} \quad \begin{array}{c} \uparrow \\ \text{size}^3 \\ \underline{1-10^9} \end{array}$$

feature scaling

Choice of features

